

Web Information Extraction and Retrieval

Programming Assignment 3:

Document indexing and querying

Marko Prelevikj	Gojko Hajduković
63130345	63180431
<code>mp2638@student.uni-lj.si</code>	<code>gh8590@student.uni-lj.si</code>
Stefan Ivanišević	
63170405	
<code>si0539@student.uni-lj.si</code>	

May 2019

1 Introduction

With the increasing development of information technologies, in the modern, digital era, most of the documents and information are stored in a digital format in order to be easy accessible to everyone within the globe. Many of the problems arose when that increasing number of data needed to be stored efficiently in order for users to have quick access to the information needed. In many information system where some type of searching is required, first natural approach in finding a related words, queries within the documents was the naive approach of sequentially looking into all of the documents for the specific words, which was very inefficient regarding to time and space. Nowadays, the most widely used and efficient way of storing the data in order to be searched from it quickly is the concept of Inverted Index.

Inverted index represents an efficient technique for storing mappings of words, content to its locations in document or in a set of documents. In this paper we introduce our implementation of first preprocessing the documents, then building an inverted index from the preprocessed content in order to allow users quick search of the content in need. We also implement the naive approach, **sequential file reading** and compare its efficiency with the approach based on inverted index. Explanation and implementation specifics are provided in following sections.

2 Data Pre-processing

For a more efficient indexing step, we needed to preprocess the corpus. The corpus contains 1416 crawled web pages. Each web page is processed in the same manner as described in continuation.

First, we extract the text data using the package [inscriptis](#). Beside that, we also normalizing the text into lowercase. Since we retrieved some of the **HTML** tags alongside with the textual data, we used regex to remove these tags.

The next step during preprocessing was to tokenize the text. We have performed tokenization (splitting up a textual data into words) using the [nltk.tokenize](#) package. During this step, we also removed all the stopwords, special characters, and the all the duplicates among the tokens.

We have also tokenized the whole textual data from the source file without removing the stop-words and special characters. We need this content in order to easily build up the snippets of the search results which are a part of the real content, and not the tokenized text.

The pre-processing output is a dictionary which contains 1416 keys denoting the source file names. The structure of our pre-processed corpus looks like this:

```
{
  "<inputFileName1>": {
    "tokens": ["<token1>", "<token2>", ... , "<tokenN>"],
    "content": ["<word1>", "<word2>", ... , "<wordN>"]
  },
  "<inputFileName2>": {
    "tokens": ["<token1>", "<token2>", ... , "<tokenN>"],
    "content": ["<word1>", "<word2>", ... , "<wordN>"]
  },
  ...,
  "<inputFileNameN>": {
    "tokens": ["<token1>", "<token2>", ... , "<tokenN>"],
    "content": ["<word1>", "<word2>", ... , "<wordN>"]
  }
}
```

3 Index building

To be able to benchmark the performance of the indices, we built two different indices: *Inverted Index* and *Sequential Index*. In the following subsections we describe how we are doing that.

3.1 Inverted Index

In our implementation of Inverted index, we used a database in order to simulate the inverted index structure. The database structure is consisted of three tables:

IndexWord consists of all the words indexed from the documents in the corpus, i.e. our dictionary.

Posting consists of a word from *IndexWord*, a document name in which the specific it appears, the frequency of appearance in the document, and the indices where the word appears in the source document.

Existing consists of a column `doesExist` that we have added in order to store a single boolean value indicating whether the inverted index has been built. We introduced this to know when we need to perform a re-initialization of the index because that is a costly operation.

In order to construct the *Inverted Index* out of a given pre-processed corpus of words for each file, we have constructed dictionary data structure. It allows us while iterating through the pre-processed corpus of words to store a word out of a given set of words as a key whose values

are number of occurrences of a word in each iterated file along with its frequency and indexes of occurrences.

Next step in our implementation is storing the constructed Inverted index in the database. First, we store the list of the keys from the *Inverted Index* dictionary which represent the set of all unique words from a corpus in the table IndexWord, then we construct a list which holds all postings in the format

```
[(word, documentName, frequency, indexes)].
```

3.2 Sequential index

The *Sequential Index* does not require a special procedure of preparing in order to perform the search. We simply use the pre-processed data to perform search on.

4 Data retrieval

The data retrieval process (search) is performed on a user provided query which is first pre-processed, to get it in the same form as the rest of the corpus, i.e. it is tokenized. Afterward we are performing the search based on which index has been chosen by the user (either sequential or inverted). Both methods return the results in the same format: a list of tuples containing the cumulative *frequency* per document, the *document* name, and the aggregated *indices* of all results.

4.1 Inverted Index

Search the *Inverted Index* is performed with a single query on the database. The query is shown in Listing ??, and it is an excerpt of the *Python* code which is performing the query.

```
SELECT documentName, sum(frequency) as freq, group_concat(indexes)
FROM Posting
— the following part is filled by Python based on
— the length of the tokenized query
WHERE word IN ({','.join(['?']*len(query))})
GROUP BY documentName
ORDER BY freq DESC
```

The query groups together the postings which match the words in the tokenized query, sums up the frequencies from the corresponding files, and aggregates the indices from the source document. The end result is a list of tuples, as previously described.

4.2 Sequential Index

The search for the *Sequential Index* is really simple: We get the pre-processed corpus and we iterate throughout the tokenized file content to get all the matches per file, and keep the track of the matches.

5 Implementation details

Some further implementation details worth noting:

Output printing Each row from the obtained result is printed in a table as provided in the [instructions](#). The printing is provided by [texttable](#).

Snippets The snippets mark the queried word with * symbols, and include up to 3 words to the left and to the right of the result.

REPL mode It is possible to enter in an interactive mode where the user is able to: *query* both types of indices, *change* the type of index being queried, force a *recreation* of the index and change the *number of results* printed in the table.

6 Query Results

This section introduces results of six different queries consisted of up to five words. We obtained identical results by querying both the inverted index and by performing a sequential search in the entire corpus of each word from the given query. The results from the three predefined queries are given in Table 2.

Querying the inverted index database is ≈ 40 times faster than sequentially searching through the corpus. We concluded this by implementing a benchmarking function which generates 10 queries consisted of 1–5 randomly selected words from the corpus (each word is independently selected), and executes the search N times. We provide the results in Table 1. This test also shows us that the more times we query the index sequentially, the faster we obtain the results. This is as a consequence of an internal caching of the results within python. Another factor is the amount of words the query is consisted of and how many times it appears within the corpus. We present three additional query results in Table 3, Table 4 and Table 5, respectively.

# iterations ¹	Inverted index	Sequential index
50	1.71ms	68.34 ms
100	1.34 ms	61.63 ms
250	1.15 ms	56.43 ms
500	1.58 ms	55.47 ms
1000	1.19 ms	56.73 ms

Table 1: Benchmark test on 10 random queries against and #iterations of both the inverted and sequential index.

7 Analysis

The inverted index has a total of 29667 words, and 1416 documents. We summed up the basic statistics of the documents in Table 1. Based on the statistics shown in Table 1 we can conclude

¹Number of iterations performed on each of the 10 randomly generated queries.

that we most retrieved the most tokens from the document named `podatki.gov.si.340.html`, which consequently leads to having the most occurrences which ranks it as the biggest document.

We summed up the basic statistics of the words occurring in our provided corpus in Table 2. We also calculated a weighted mean frequency of the words from the corpus with the formula presented in Equation 1.

$$f_{wi} = \frac{\sum_{d \in D} f_{i,d}}{\sum_{d \in D} f_d} * \frac{\sum_{d \in D} occurrence_d(i)}{|D|} \quad (1)$$

where f_{wi} denotes the weighted frequency of the i^{th} word, $f_{i,d}$ denotes the frequency of the i^{th} word in the d^{th} document, n_d denotes the sum of frequencies of all words appearing in the d^{th} word. The function $occurrence_d(i)$ is calculated as presented in Equation 2. D denotes the set of documents, i.e. the corpus, whereas $|D|$ denotes the number of documents in the corpus. Translated into *SQL*, we used the query in the listing below.

$$occurrence_d(i) = \begin{cases} 1 & i \in d \\ 0 & i \notin d \end{cases} \quad (2)$$

```
SELECT word, avg(frequency) * sum(1)/1416.0 weightedFreq, sum(1) occurrence
FROM "main"."Posting"
GROUP BY word
ORDER BY weightedFreq DESC
```

We went one more step ahead, and we calculated the distribution of the average frequency of words occurring within the document. The visualization of our results is shown in Figure 1. The visualized distribution looks like it follows the Poisson distribution, which can be expected and interpreted as having websites with an average number of words occurring on it, and each time we visit one we expect to have so many words, but then there are some documents, such as `podatki.gov.si.340.html`, which consists of a lot more words because it is some kind of a dictionary, and it contains a lot of meaningful words at one place.

8 Conclusion

Throughout the paper we introduced our implementation of building inverted index and query against it. Key differences between naive approach of sequential file reading and inverted index are noted supporting with experiments and results which prove how much more efficient the inverted index approach is, which is the reason why that concept is widely used nowadays even though Inverted Index takes much more time in constructing the structure as it allows very efficient querying.

²See Section 7 for more details

³*#Occurrences denotes in how many documents the word appears.




Figure 1: Distribution of the average frequency of tokens per document.

Query	Rank	Frequency	Document	Snippet
predelavne dejavnosti	1	75	evem.gov.si.377	defektolog v zdravstveni *dejavnosti* + dekan ozirom.. dietetik v zdravstveni *dejavnosti* + dimnikar + ...
	2	40	podatki.gov.si.340	nosilec dopolnilne *dejavnosti* na kmetiji bregar ... šport center interesnih *dejavnosti* ptuj center judovske ... šolskih in obšolskih *dejavnosti* center urbane kulture ...
	3	36	evem.gov.si.452	evem>dejavnosti>druge storitvene *dejavnosti* , druge nerazvrščene ...)druge storitvene *dejavnosti* , druge nerazvrščene ... skd šifra zajema *dejavnosti* in storitve , ... začetek in opravljanje *dejavnosti* .predpisi in...
trgovina	1	95	evem.gov.si.651	govedoreja + druga *trgovina* na drobno v ... prodajalnah + druga *trgovina* na drobno v ...prodajalnah + druga *trgovina* na drobno v ... živili + druga *trgovina* na drobno zunaj ...
	2	92	evem.gov.si.21	moj e-vem evem>področja *trgovina* tu boste našli ... dejavnosti * druga *trgovina* na drobno v ... prodajalnah * druga *trgovina* na drobno zunaj ...) * nespecializirana *trgovina* na debelo * ...
	3	82	podatki.gov.si.340	a dent , *trgovina* in storitve , adria investicije *trgovina* , posredništvo , ... d.o.o . ahatservis *trgovina* in storitve , ... d.o.o . alba *trgovina* in proizvodnja , ...
social services	1	5	e-uprava.gov.si.45	, retirement * *social* services , health ... retirement * social *services* , health , ... relationship etc. ? *social* services , health ... etc. ? social *services* , health , ... i obtain financial *social* assistance ? how ...
	2	5	e-uprava.gov.si.5	, retirement * *social* services , health ... retirement * social *services* , health , ... relationship etc. ? *social* services , health ... etc. ? social *services* , health , ... i obtain financial *social* assistance ? how ...
	3	1	evem.gov.si.661	records and related *services* (ajpes) ...

Table 2: Results for pre-defined queries

Rank	Frequency	Document	Snippet
1	62	evem.gov.si.661	english welcome to *the* e-trr system ! ...) account for *the* payment of the ... the payment of *the* initial capital of ... a precondition for *the* establishing of a ... a company in *the* register of companies ...
2	51	podatki.gov.si.490	povezava header - *middle* left navigation * ... + kmetijstvo , *ribištvo* , gozdarstvo in ... , ribištvo , *gozdarstvo* in prehrana + ... portalu header - *middle* right navigation * ... povezave header - *middle* left navigation * ...
3	49	podatki.gov.si.239	r povezava header - *middle* left navigation * ... + kmetijstvo , *ribištvo* , gozdarstvo in ... , ribištvo , *gozdarstvo* in prehrana + ... portalu header - *middle* right navigation * ... povezave header - *middle* left navigation * ...

Table 3: Results of the query "the middle gozdarstvo kontakt ribištvo"

Rank	Frequency	Document	Snippet
1	10	podatki.gov.si.494	državni organi energetika *cene* električne energije za ... z naslovom " *cene* električne ... nadaljujte ... državni organi energetika *cene* energentov , slovenija ... z naslovom " *cene* energentov , državni organi energetika *cene* za zemeljski plin ...
2	6	podatki.gov.si.506	državni organi energetika *cene* za zemeljski plin ... z naslovom " *cene* za zemeljski plin ... državni organi energetika *cene* za zemeljski plin ... z naslovom " *cene* za zemeljski plin ... državni organi energetika *cene* električne energije za ...
3	6	podatki.gov.si.508	državni organi energetika *cene* električne energije za ... z naslovom " *cene* električne ... nadaljujte ... državni organi energetika *cene* za zemeljski plin ... z naslovom " *cene* za zemeljski plin ... državni organi energetika *cene* za zemeljski plin ...

Table 4: Results of the query "cene 2018"

Rank	Frequency	Document	Snippet
1	13	evem.gov.si.227	pogoji : * *zakon* o sodiščih in ... in če drug *zakon* ali uredba določa ... za pravosodje * *zakon* o državnem tožilstvu ... jih določa ta *zakon* , ter opravlja ... za pravosodje * *zakon* o alternativnem reševanju ...
2	11	evem.gov.si.224	pogoji : * *zakon* o varstvu okolja ... in prostor * *zakon* o zasebnem varovanju ... veljavno licenco . *zakon* o zasebnem varovanju ... za obrambo * *zakon* o varnosti v ... za infrastrukturo * *zakon* o varstvu pred ...
3	9	e-prostor.gov.si.13	” . * *zakon* o evidentiranju nepremičnin ...) . * *zakon* o nadzoru države ... 35/2010) ta *zakon* določa organizacijo in ... v 10. členu *zakon* določa , da ... v 12. členu *zakon* določa pristojnost geodetske ...

Table 5: Results of the query "naslovni zakon"

Page	#Tokens	#Occurrences	Mean #Occurrences
podatki.gov.si.340	6528	27421	4.201
e-prostor.gov.si.57	1679	3591	2.139
evem.gov.si.398	1559	4252	2.727
evem.gov.si.651	1292	2598	2.011
e-uprava.gov.si.56	1191	2293	1.925

Table 6: Documents with most words

Word	Frequency	Weighted Mean Frequency ²	#Occurrences ³
podatkov	11000	7.768	863
slovenije	8814	6.225	896
republike	8356	5.901	1165
podatki	4931	3.482	732
navigation	4474	3.160	561

Table 7: Basic statistics of the words from the corpus.